

EngTrace inference pricing

Inference means each model solves every EngTrace problem once. It covers generating the answers only, not evaluating them.

EngTrace has five engineering branches: chemical, civil, electrical, industrial and mechanical. Each branch has 30 problem templates, and each template is used 15 times with different values. That gives $5 \times 30 \times 15 = 2,250$ problems (870 easy, 870 intermediate and 510 advanced), so each model makes 2,250 runs.

Open-source models

Through OpenRouter, cheapest upstream serving fp8 or better.

Model	Input \$/M	Output \$/M	Est. cost
openai/gpt-oss-20b	0.018	0.090	\$1.07
google/gemma-4-26b-a4b-it	0.090	0.300	\$3.59
deepseek/deepseek-v4.1-flash	0.140	0.420	\$5.03
qwen/qwen3-235b-a22b	0.087	0.350	\$4.17
z-ai/glm-5.3-flash	0.075	0.250	\$2.99
z-ai/glm-5.3	1.400	4.400	\$52.65
meta/muse-glimmer-30b	0.300	1.100	\$13.13
moonshotai/kimi-k3	2.100	10.950	\$130.18
Total			\$212.81

Frontier models

Model	Input \$/M	Output \$/M	Est. cost
openai/gpt-5.4-nano	0.20	1.25	\$14.84
openai/gpt-5.4-mini	0.75	4.50	\$53.43
google/gemini-3.1-flash-lite	0.25	1.50	\$17.81
anthropic/claude-sonnet-5	2.00	10.00	\$118.94
Total			\$205.01