

EngTrace inference pricing

Inference is each model solving every EngTrace problem once; evaluating the answers is not included. EngTrace covers five engineering branches (chemical, civil, electrical, industrial and mechanical), each with 30 problem templates. Every template is instantiated 15 times with different sampled values, giving $5 \times 30 \times 15 = 2,250$ problems (870 easy, 870 intermediate, 510 advanced), so each model makes 2,250 runs. The estimated cost is the total over those runs, assuming each model writes as much as GPT-5 did in EngTrace's evaluator pilot: about 270 input and 5,232 output tokens per problem, reasoning included.

Open-source models

Through OpenRouter, cheapest upstream serving fp8 or better.

Model	Input \$/M	Output \$/M	Est. cost
openai/gpt-oss-20b	0.018	0.090	\$1.07
google/gemma-4-26b-a4b-it	0.090	0.300	\$3.59
deepseek/deepseek-v4.1-flash	0.140	0.420	\$5.03
qwen/qwen3-235b-a22b	0.087	0.350	\$4.17
z-ai/glm-5.3-flash	0.075	0.250	\$2.99
z-ai/glm-5.3	1.400	4.400	\$52.65
meta/muse-glimmer-30b	0.300	1.100	\$13.13
moonshotai/kimi-k3	2.100	10.950	\$130.18
Total			\$212.81

Frontier models

Model	Input \$/M	Output \$/M	Est. cost
openai/gpt-5.4-nano	0.20	1.25	\$14.84
openai/gpt-5.4-mini	0.75	4.50	\$53.43
google/gemini-3.1-flash-lite	0.25	1.50	\$17.81
anthropic/claude-haiku-4.5	1.00	5.00	\$59.47
anthropic/claude-sonnet-5	2.00	10.00	\$118.94
Total			\$264.48