

# EngTrace inference pricing

Inference is each model solving every EngTrace problem once: 150 templates × 15 instances = 2,250 runs per model.

## Open-source models

*Through OpenRouter, cheapest upstream serving fp8 or better.*

| Model                        | Input \$/M | Output \$/M | Est. cost     |
|------------------------------|------------|-------------|---------------|
| openai/gpt-oss-20b           | 0.018      | 0.090       | <b>\$1.07</b> |
| google/gemma-4-26b-a4b-it    | 0.090      | 0.300       | <b>\$3.59</b> |
| deepseek/deepseek-v4.1-flash | 0.140      | 0.420       | <b>\$5.03</b> |
| qwen/qwen3-235b-a22b         | 0.087      | 0.350       | <b>\$4.17</b> |
| z-ai/glm-5.3-flash           | 0.075      | 0.250       | <b>\$2.99</b> |
| z-ai/glm-5.3                 | 1.400      | 4.400       | <b>\$53</b>   |
| meta/muse-glimmer-30b        | 0.300      | 1.100       | <b>\$13</b>   |
| moonshotai/kimi-k3           | 2.100      | 10.950      | <b>\$130</b>  |
| <b>Total</b>                 |            |             | <b>\$213</b>  |

## Frontier models

| Model                        | Input \$/M | Output \$/M | Est. cost    |
|------------------------------|------------|-------------|--------------|
| openai/gpt-5.4-nano          | 0.20       | 1.25        | <b>\$15</b>  |
| openai/gpt-5.4-mini          | 0.75       | 4.50        | <b>\$53</b>  |
| google/gemini-3.1-flash-lite | 0.25       | 1.50        | <b>\$18</b>  |
| anthropic/claude-haiku-4.5   | 1.00       | 5.00        | <b>\$59</b>  |
| anthropic/claude-sonnet-5    | 2.00       | 10.00       | <b>\$119</b> |
| <b>Total</b>                 |            |             | <b>\$264</b> |