

EngTrace inference pricing

Model prices per million tokens and the estimated cost of one full benchmark run

Prices read on OpenRouter, 19–21 September 2026

The tables list each model’s input and output price. The last two columns estimate what **one full EngTrace run** costs for that model: 150 templates × 15 instances = 2,250 items, inference only, no evaluation. They use token counts measured on EngTrace’s own evaluator pilot, because a model’s cost depends on how much it writes, and that is not known in advance for a new model:

- **Direct answer:** about 318 input and 813 output tokens per item, the average of the pilot’s non-reasoning models.
- **Reasons like GPT-5:** about 270 input and 5,232 output tokens per item, GPT-5’s measured average including its reasoning.

A model’s real cost will usually fall between the two. The heaviest reasoner in the pilot, DeepSeek-R1, wrote about 10,000 output tokens per item, roughly twice the second column.

Two things to watch

- **Output is where the money goes for reasoning models.** Output prices are four to eight times input, and engineering problems produce long reasoning, so the cost tracks how much each model writes.
- **OpenRouter prices move between upstreams from day to day.** These were read around 19 to 21 September 2026 and should be re-read before the run.

Open-weight models, through OpenRouter

Cheapest upstream serving fp8 or better.

Model	Input \$/M	Output \$/M	Direct answer	Reasons like GPT-5
openai/gpt-oss-20b	0.018	0.090	\$0.18	\$1.07
google/gemma-4-26b-a4b-it	0.090	0.300	\$0.61	\$3.59
google/gemma-4-31b-it	0.130	0.380	\$0.79	\$4.55
deepseek/deepseek-v4-flash	0.036	0.071	\$0.16	\$0.86
deepseek/deepseek-v4.1-flash	0.140	0.420	\$0.87	\$5.03
qwen/qwen3-235b-a22b	0.087	0.350	\$0.70	\$4.17
qwen/qwen3.8-27b	0.150	1.875	\$3.54	\$22
z-ai/glm-5.3-flash	0.075	0.250	\$0.51	\$2.99
z-ai/glm-5.3	1.400	4.400	\$9.05	\$53
meta/muse-glimmer-30b	0.300	1.100	\$2.23	\$13
moonshotai/kimi-k3	2.100	10.950	\$22	\$130

OpenAI

Model	Input \$/M	Output \$/M	Direct answer	Reasons like GPT-5
gpt-5-nano †	0.05	0.40	\$0.77	\$4.74
gpt-5.6-luna	0.20	1.20	\$2.34	\$14
gpt-5.4-nano	0.20	1.25	\$2.43	\$15
gpt-5-mini †	0.25	2.00	\$3.84	\$24
gpt-5.4-mini	0.75	4.50	\$8.77	\$53
gpt-5 †	1.25	10.00	\$19	\$118
gpt-5.6-terra	2.00	12.00	\$23	\$142
gpt-5.6-sol	4.00	20.00	\$39	\$238
gpt-5.5	5.00	30.00	\$58	\$356
gpt-6-astra	10.00	50.00	\$99	\$595

Google

Model	Input \$/M	Output \$/M	Direct answer	Reasons like GPT-5
gemini-2.5-flash-lite †	0.10	0.40	\$0.80	\$4.77
gemini-3.1-flash-lite	0.25	1.50	\$2.92	\$18
gemini-2.5-flash †	0.30	2.50	\$4.79	\$30
gemini-3.6 / 3.7 / 3.8-flash	0.75	3.75	\$7.40	\$45
gemini-2.5-pro †	1.25	10.00	\$19	\$118
gemini-3.5-flash	1.50	9.00	\$18	\$107
gemini-3.1-pro-preview	2.00	12.00	\$23	\$142

Anthropic

Model	Input \$/M	Output \$/M	Direct answer	Reasons like GPT-5
claude-haiku-4.5	1.00	5.00	\$9.86	\$59
claude-sonnet-5	2.00	10.00	\$20	\$119
claude-opus-5	5.00	25.00	\$49	\$297
claude-fable-5-1	10.00	50.00	\$99	\$595

† Retires before publication: gpt-5, gpt-5-mini and gpt-5-nano on 11 December 2026; the Gemini 2.5 models from 20 October 2026.

How the estimates are made

- $\text{Cost} = 2,250 \times (\text{input tokens} \times \text{input price} + \text{output tokens} \times \text{output price}) \div 1,000,000$.
- Token counts come from the pilot's answers to 60 items covering all five branches and all three difficulty levels, under the benchmark's own prompt.
- Re-runs after truncated or failed answers are not included; the pilot needed 0–15% extra for most models and up to 45% for a model whose reasoning ran away.
- Half-price batch endpoints, where a provider offers them, would roughly halve the frontier rows.