

EngTrace inference pricing

Inference is each model solving every EngTrace problem once: 150 templates × 15 instances = 2,250 runs per model.

Open-weight models, through OpenRouter

Cheapest upstream serving fp8 or better.

Model	Input \$/M	Output \$/M	Est. cost
openai/gpt-oss-20b	0.018	0.090	\$1.07
google/gemma-4-26b-a4b-it	0.090	0.300	\$3.59
google/gemma-4-31b-it	0.130	0.380	\$4.55
deepseek/deepseek-v4-flash	0.036	0.071	\$0.86
deepseek/deepseek-v4.1-flash	0.140	0.420	\$5.03
qwen/qwen3-235b-a22b	0.087	0.350	\$4.17
qwen/qwen3.8-27b	0.150	1.875	\$22
z-ai/glm-5.3-flash	0.075	0.250	\$2.99
z-ai/glm-5.3	1.400	4.400	\$53
meta/muse-glimmer-30b	0.300	1.100	\$13
moonshotai/kimi-k3	2.100	10.950	\$130

OpenAI

Model	Input \$/M	Output \$/M	Est. cost
gpt-5-nano †	0.05	0.40	\$4.74
gpt-5.6-luna	0.20	1.20	\$14
gpt-5.4-nano	0.20	1.25	\$15
gpt-5-mini †	0.25	2.00	\$24
gpt-5.4-mini	0.75	4.50	\$53
gpt-5 †	1.25	10.00	\$118
gpt-5.6-terra	2.00	12.00	\$142
gpt-5.6-sol	4.00	20.00	\$238
gpt-5.5	5.00	30.00	\$356
gpt-6-astra	10.00	50.00	\$595

Google

Model	Input \$/M	Output \$/M	Est. cost
gemini-2.5-flash-lite †	0.10	0.40	\$4.77
gemini-3.1-flash-lite	0.25	1.50	\$18
gemini-2.5-flash †	0.30	2.50	\$30
gemini-3.6 / 3.7 / 3.8-flash	0.75	3.75	\$45
gemini-2.5-pro †	1.25	10.00	\$118
gemini-3.5-flash	1.50	9.00	\$107
gemini-3.1-pro-preview	2.00	12.00	\$142

Anthropic

Model	Input \$/M	Output \$/M	Est. cost
claude-haiku-4.5	1.00	5.00	\$59
claude-sonnet-5	2.00	10.00	\$119
claude-opus-5	5.00	25.00	\$297
claude-fable-5-1	10.00	50.00	\$595

† Retires before publication: gpt-5, gpt-5-mini and gpt-5-nano on 11 December 2026; the Gemini 2.5 models from 20 October 2026.

How the estimates are made

- Est. cost = $2,250 \times (\text{input tokens} \times \text{input price} + \text{output tokens} \times \text{output price}) \div 1,000,000$, the full cost of all 2,250 runs.
- It assumes each model writes as much as GPT-5 did in EngTrace's evaluator pilot: about 270 input and 5,232 output tokens per problem, reasoning included. A model that answers without reasoning writes about 800 output tokens and would cost roughly a sixth of the figure shown.
- Not included: re-runs of truncated or failed answers (0–15% extra for most models in the pilot) and half-price batch endpoints.
- Prices as read on OpenRouter, 19–21 September 2026; they move between upstreams from day to day.